

SCHOOL OF DATA ANALYSIS

ML in HEP: Aspects and Applications

Tatiana Likhomanenko

Introduction




ML in HEP

- › High-level triggers
- › Calibration (energy in the calorimeter)
- › Particle identification
- › Stripping line (rarely)
- › Tagging (jet-tagging, B-tagging)
- › Analysis (rare decays, ...)

LHC Event

- › Sample: one proton-proton bunches collision
- › Event consists of:
 - tracks (track description)
 - secondary vertices (SV description)
- › Questions:
 - How to describe event in ML terms?
 - How to train model on such samples?

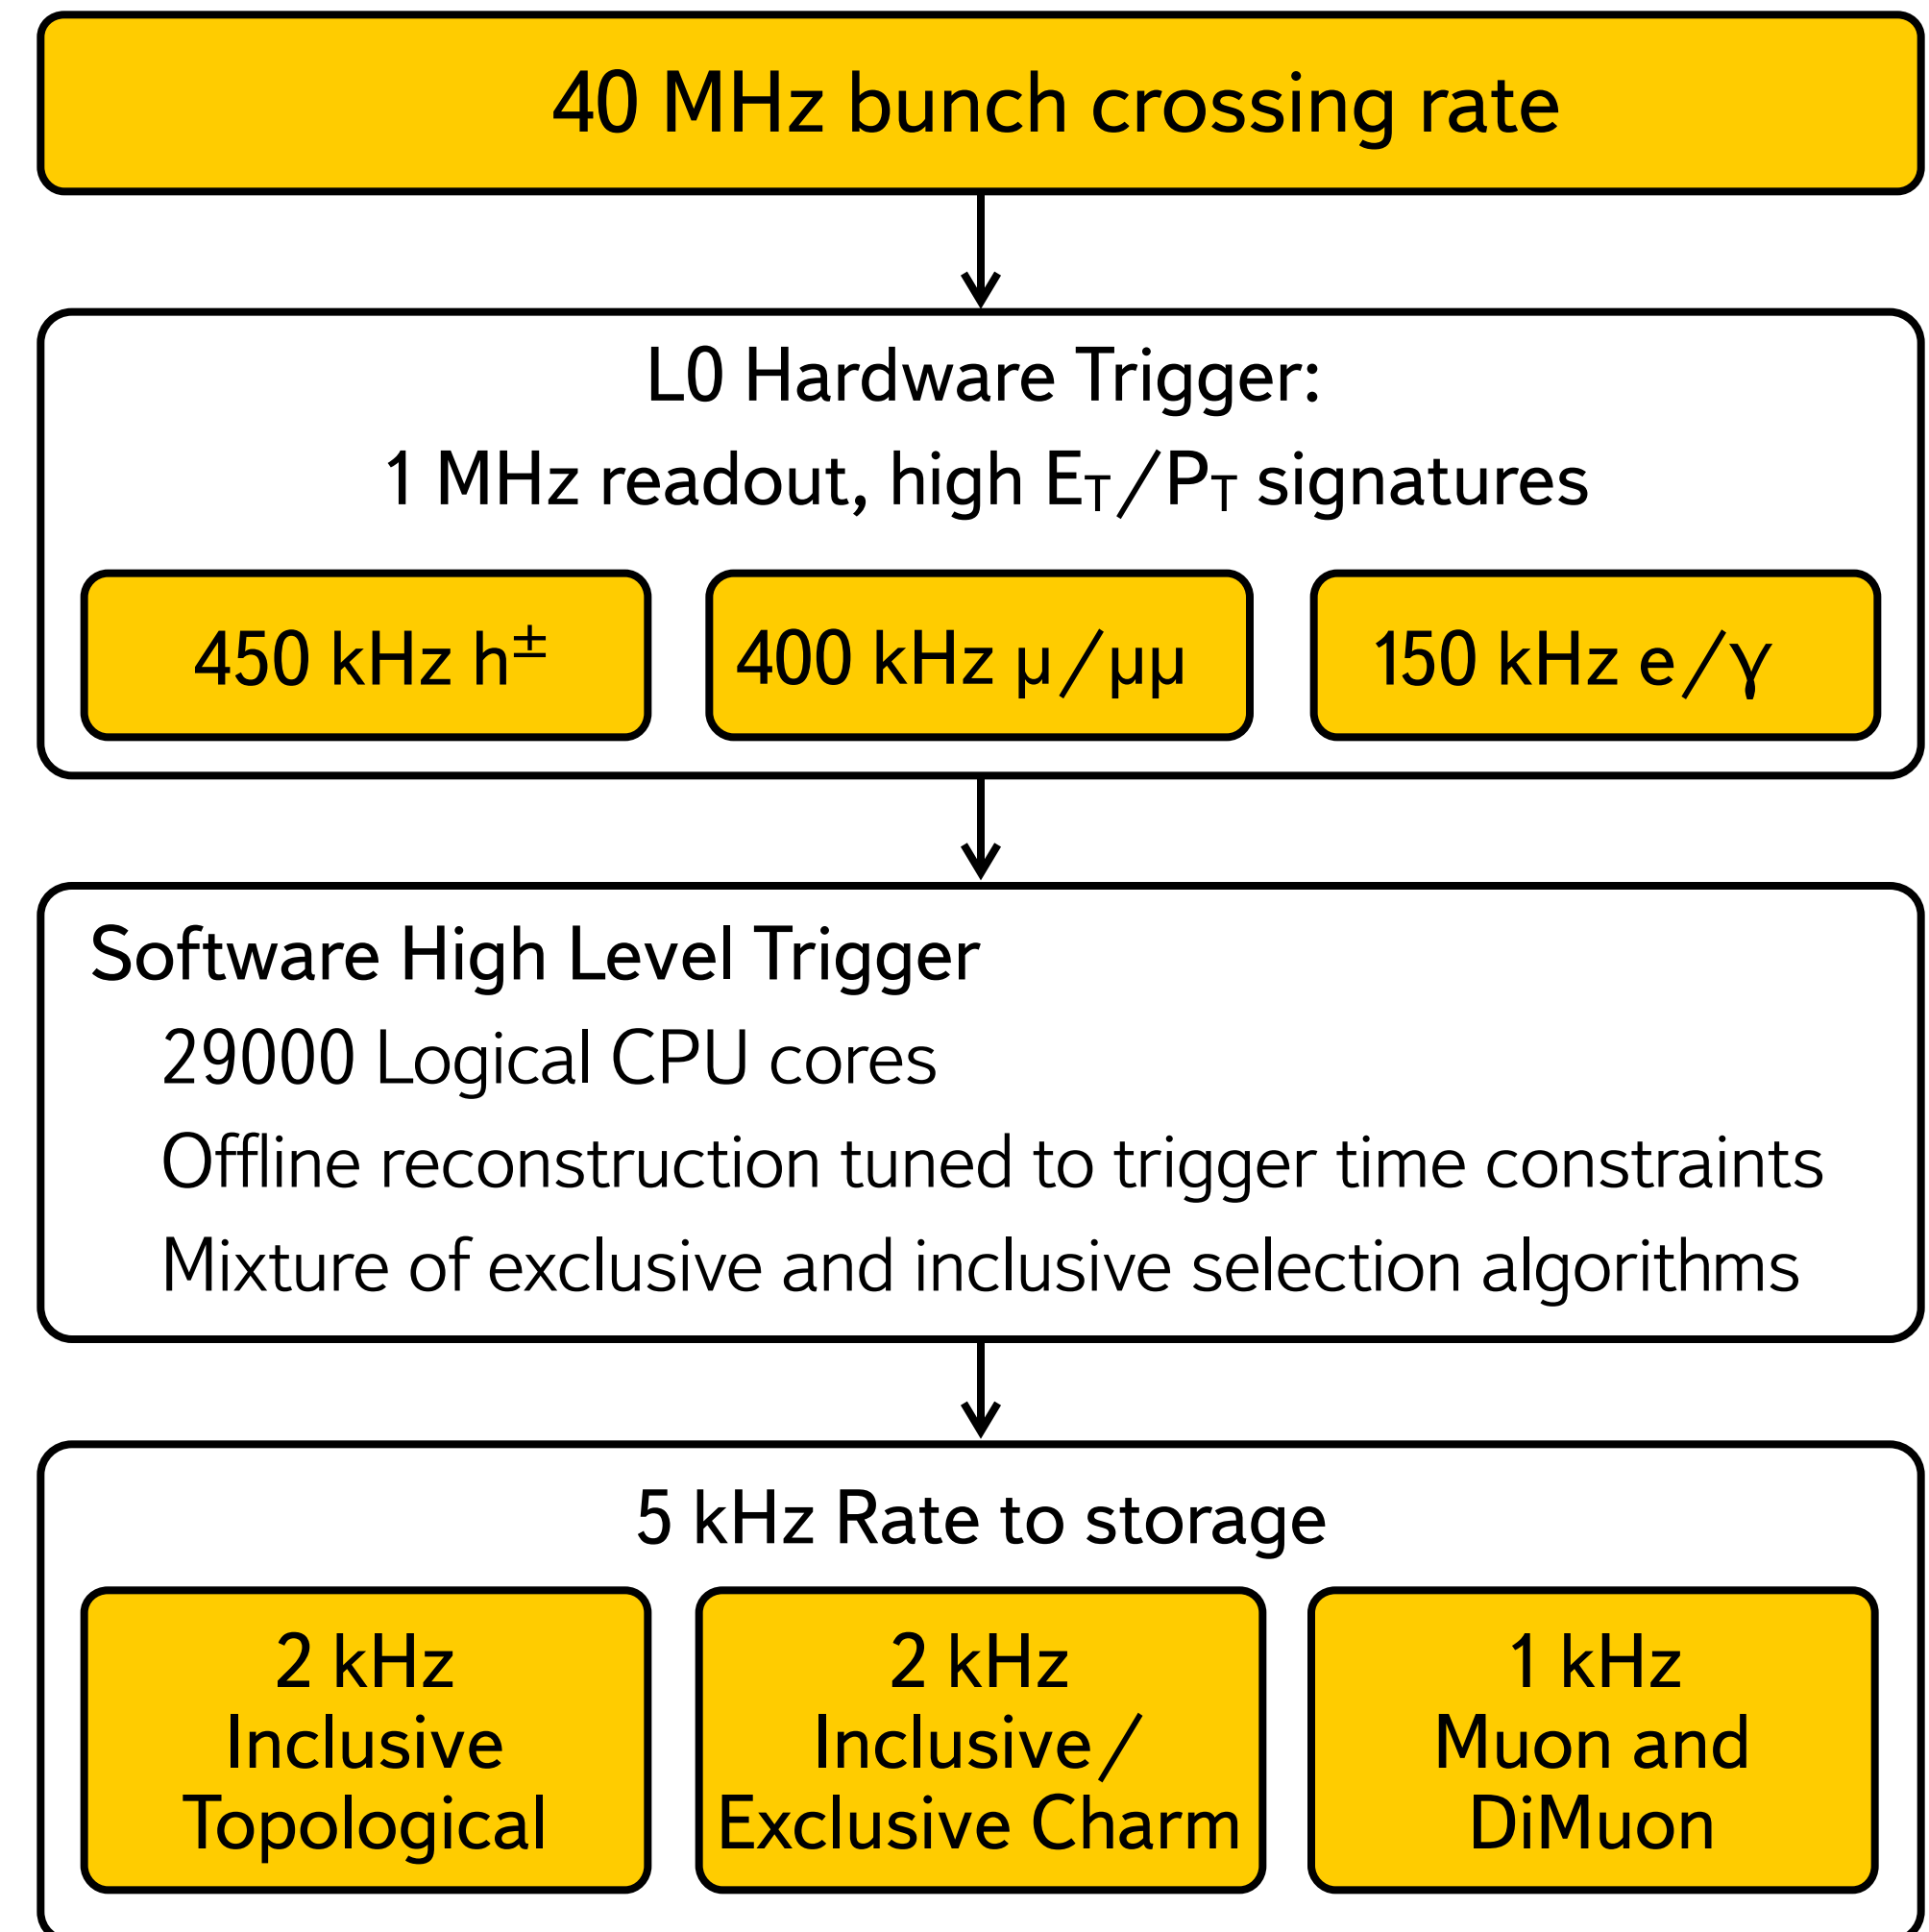


Topological trigger



LHCb trigger system

- › Select events to store them for offline processing
- › Should efficiently select interesting events
- › An event is interesting if it contains at least one interesting SV
- › Output rate for trigger system is limited

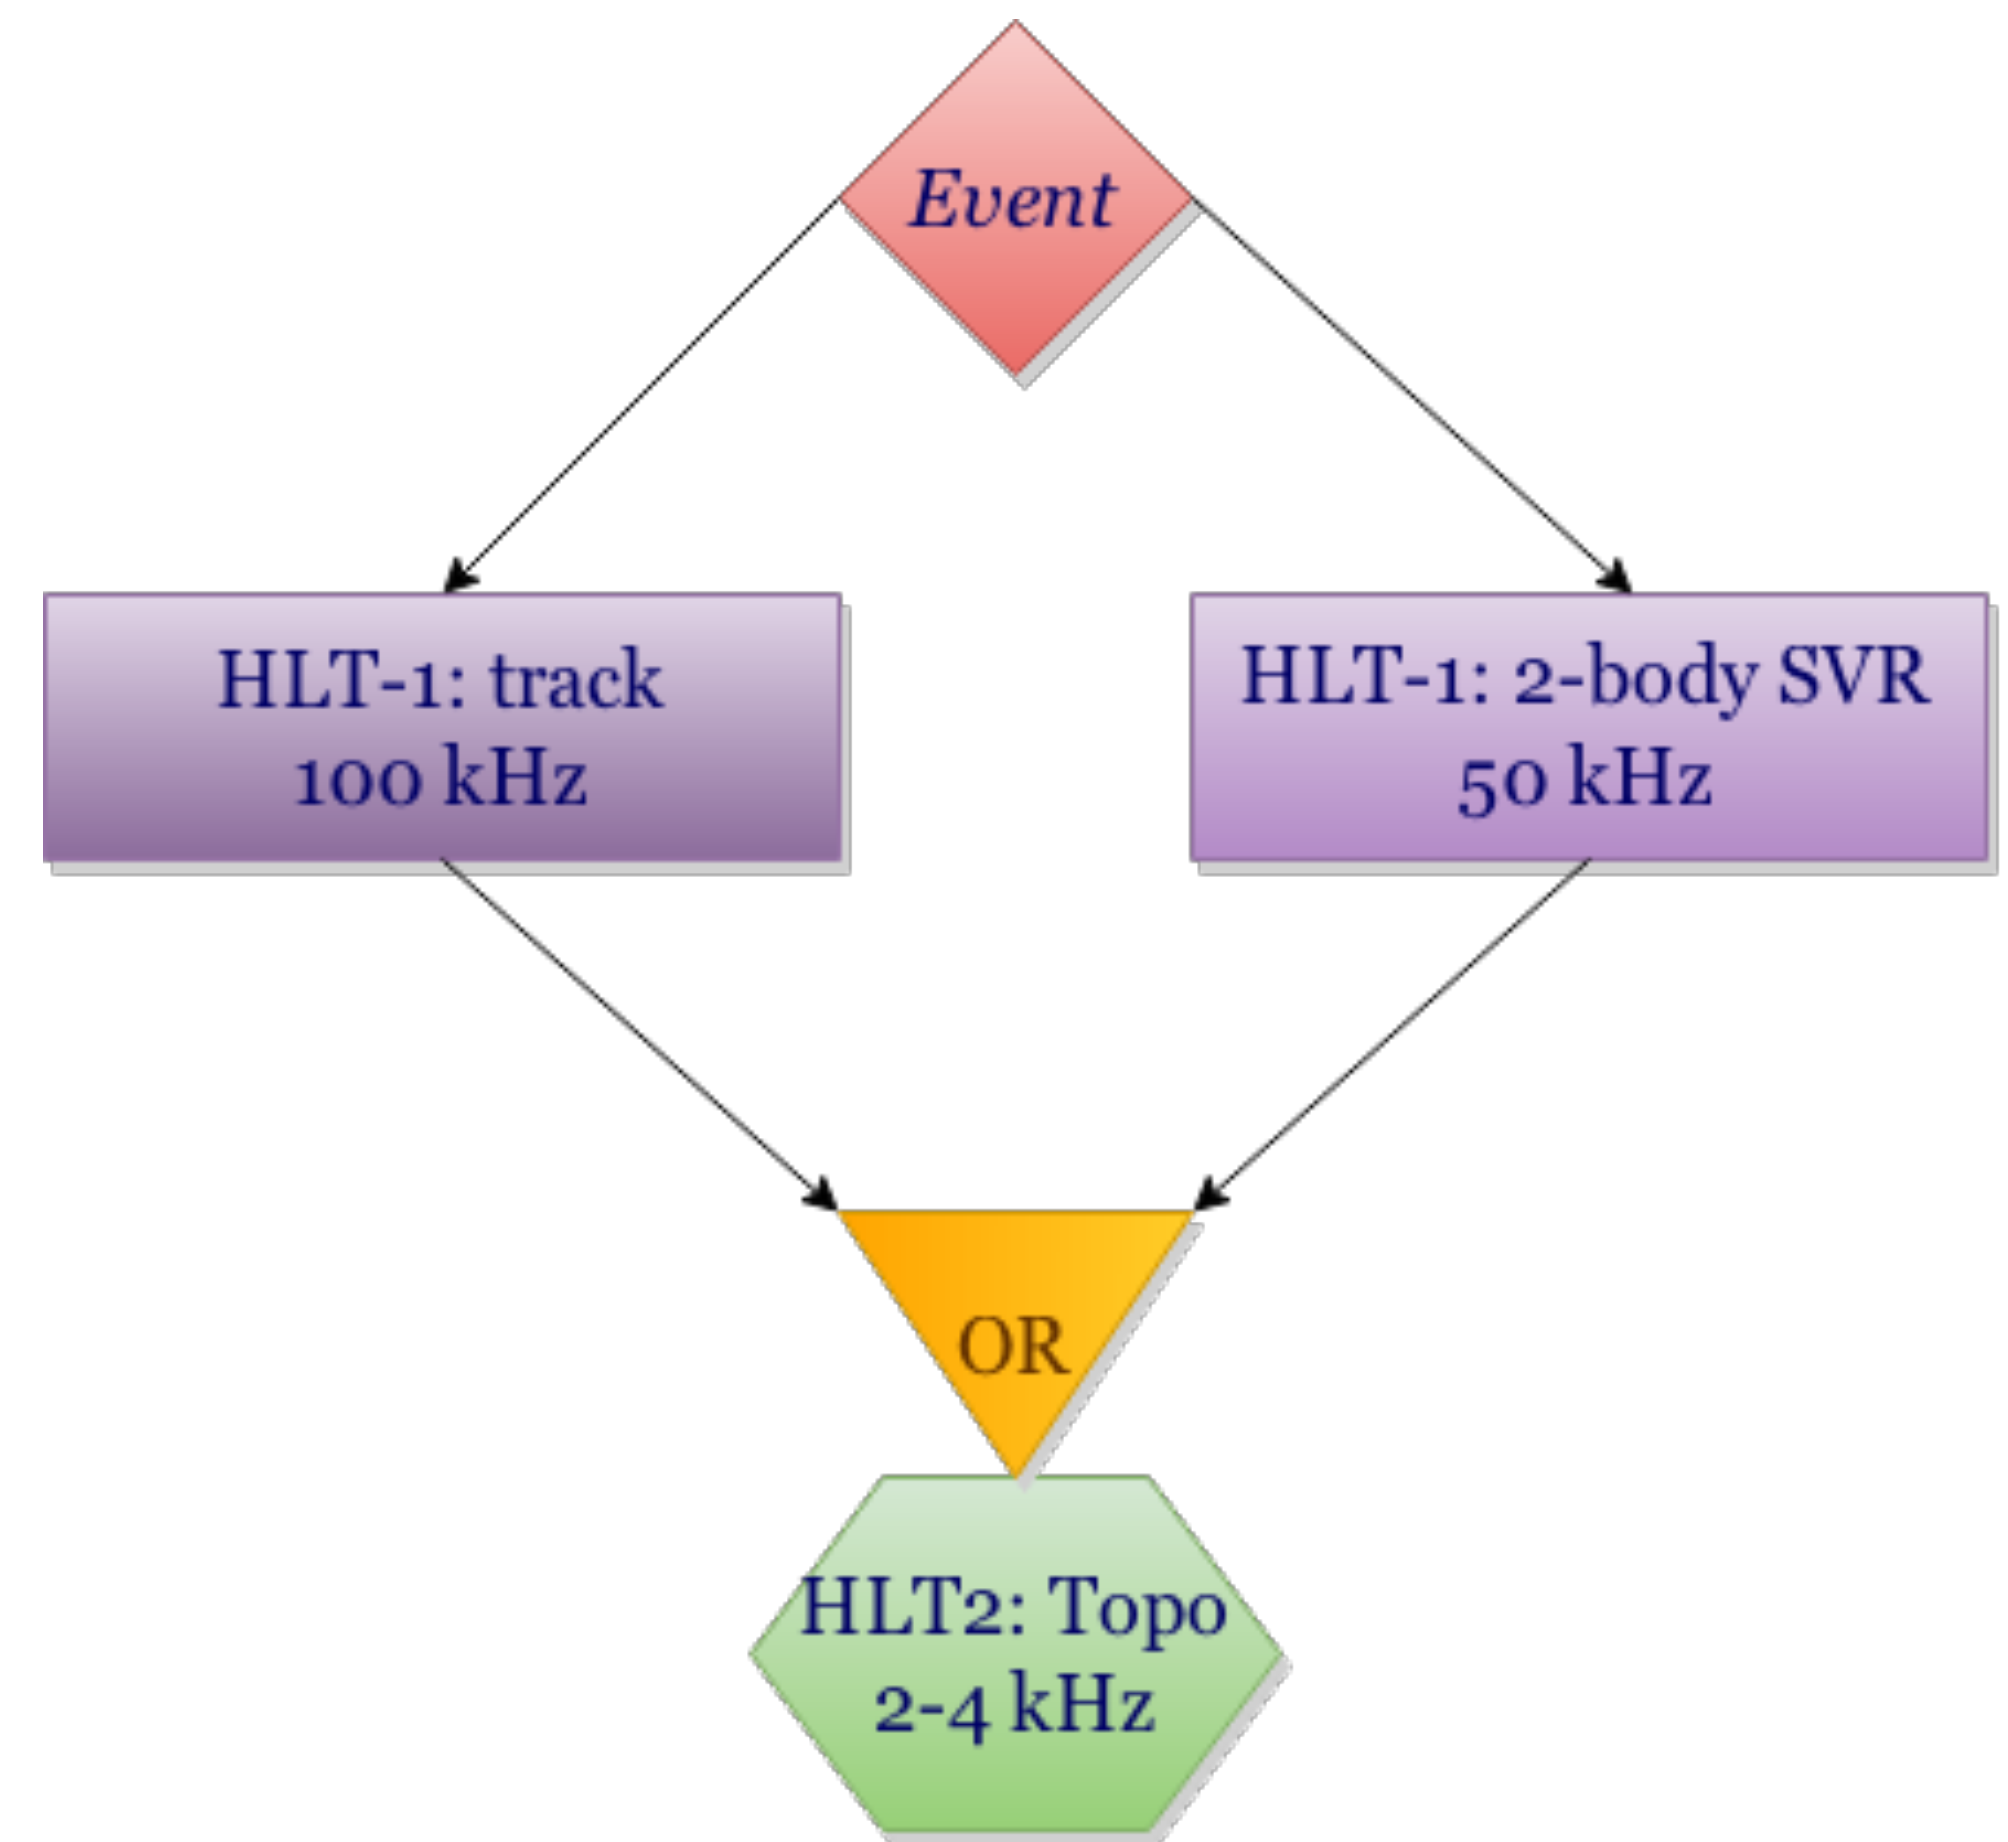


LHCb topological trigger

- › Generic trigger for decays of beauty and charm hadrons
- › It designed to be inclusive trigger line to efficiently select any B decay with at least 2 charged daughters
- › Look for 2, 3, 4 track combinations in a wide mass range
- › Designed to efficiently select decays with missing particles
- › Use fast-track fit to improve signal efficiency and minbias rejection


Run-II LHCb topological trigger

- › HLT-1 track is looking for either one super high PT or high displacement track
- › HLT-1 2-body SV classifier is looking for two tracks making a vertex
- › HLT-2 improved topo classifier uses full reconstructed event to look for 2, 3, 4 and more tracks making a vertex



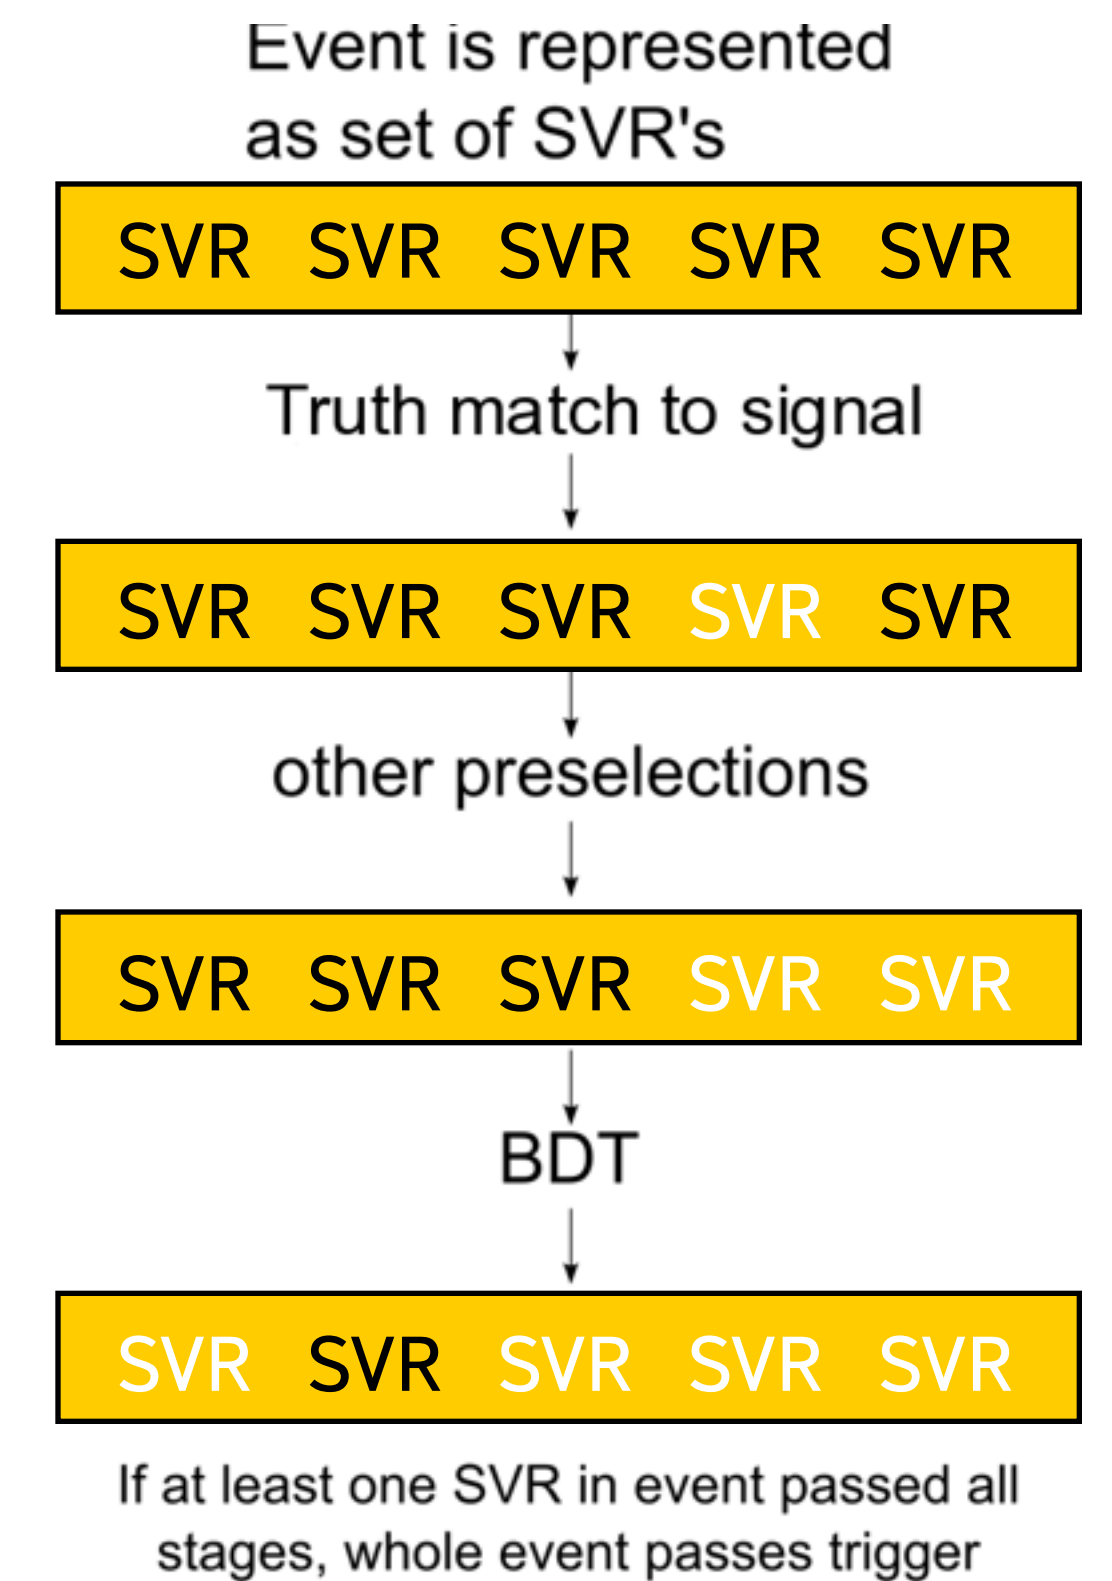
N body tracks

- › Two, three or four tracks are combined to form a SVR
- › Each secondary vertex in Monte Carlo data is preselected in such way, that all tracks must be matched to particles from the signal decay (true match preselection)




Data

- › Monte Carlo samples (used as signal-like) are simulated 13-TeV B decays of various topologies
- › Generic Pythia 13-TeV proton-proton collisions are used as background-like sample
- › Training data are set of SVs for all events
- › Most events have many secondary vertices (not all events have them)
- › Goal is to improve efficiency for each type of signal events along fixed efficiency for background




Event representation

How to measure quality?




ROC curve, computed for events

- › Output rate = false positive rate (FPR) for events
- › Optimize true positive rate (TPR) for fixed FPR for events
- › Weight signal events in such way that channels have the same sum of weights.
- › Optimize ROC curve in a small FPR region




ROC curve interpretation


Dependence on output rate



Hierarchical training

- › Train separate models for:
 - each channel
 - each n-body type: 2, 3, 4
- › Use them as additional features later
- › Use folding scheme to train additional features or to apply additional classifier selections




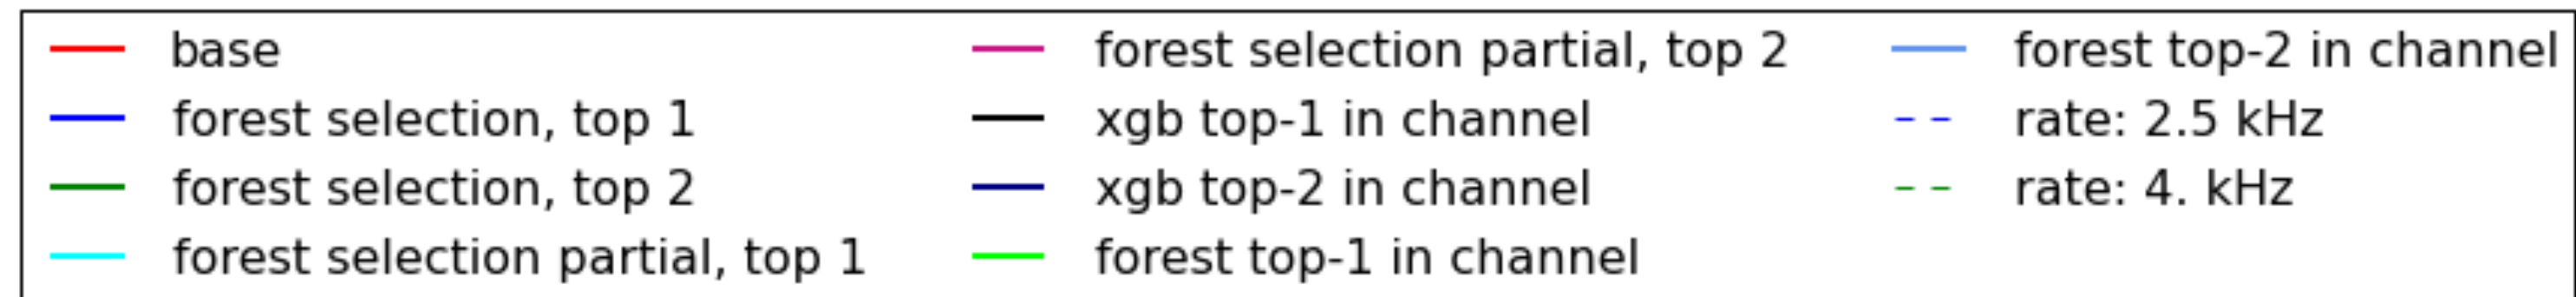


Simulated signal event
contains at least one
interesting *SV*, but not each
SV should be interesting

Random forest for SVs selection

- › Train random forest (RF) on SVs using folding scheme
 - RF is stable to noise in data
 - RF doesn't penalize in case of misclassification (can find noisy samples)
- › Select top-1, top-2 SVs by RF predictions for each signal event
- › Train classifier on selected SVs
- › Try another algorithm instead of RF, maybe it will work!

Random forest for SVs selection




Online processing

There are two possibilities to speed up prediction operation:


- › Bonsai boosted decision tree format (BBDT)
- › Post-pruning

BBDT

- › Features hashing using bins before training
- › Converting decision trees to n-dimensional table (lookup table)
- › Table size is limited in RAM (1Gb), thus count of bins for each features should be small (5 bins for each of 12 features)
- › Discretization reduces the quality
- › Prediction operation takes one reading from the table



BBDT, results




Post-pruning

- › Train GB over oblivious trees with several thousands trees
- › Reduce this amount of trees to a hundred
- › Greedily choose trees in a sequence from the initial ensemble to minimize a modified loss function:


$$\sum_{\text{signal}} \log \left(1 + e^{-F(x)} \right) + \sum_{\text{background}} e^{F(x)}$$

- › At the same time change values in leaves (tree structure is preserved)

Post-pruning, results



Topological trigger results (without RF trick)



<https://github.com/yandexdataschool/LHCb-topo-trigger>



sPlot technique





Why using it?

- › Monte Carlo is not-well simulated
- › Need to work with real unlabeled data
- › Need somehow to label real data: want to restore for features their distributions for the signal and background data
- › Our main knowledge is the mass distribution for real data from which we can extract the mass pdfs for signal and background.
- › How to restore signal/bck pdfs for other features?

Feature initial distributions



Two mass bins





$$Bin1 : w_{b_1} f_b + w_{s_1} f_s$$

$$Bin2 : w_{b_2} f_b + w_{s_2} f_s$$

$$*w_{b_2} + \text{will obtain initial signal distribution}$$

$$*(-w_{b_1})$$


Reconstruction



More bins: sWeight

- › Equivalent to some optimization problem
- › Has simple explicit solution
- › Produce weights (sWeight) for each event
- › Feature pdf with $\text{sWeight} = \text{signal pdf}$
- › Details for sPlot technique
- › Blogpost about sPlot (simplified explanation)

Proportion of events inside bins



How to reweight?

ML on sPlot data



Properties and problems

- › There are always negative values among sWeights
- › The correspondence between probability to be signal/bck and sWeight
- › Standard ML algorithms cannot be applied* (loss function is not convex)

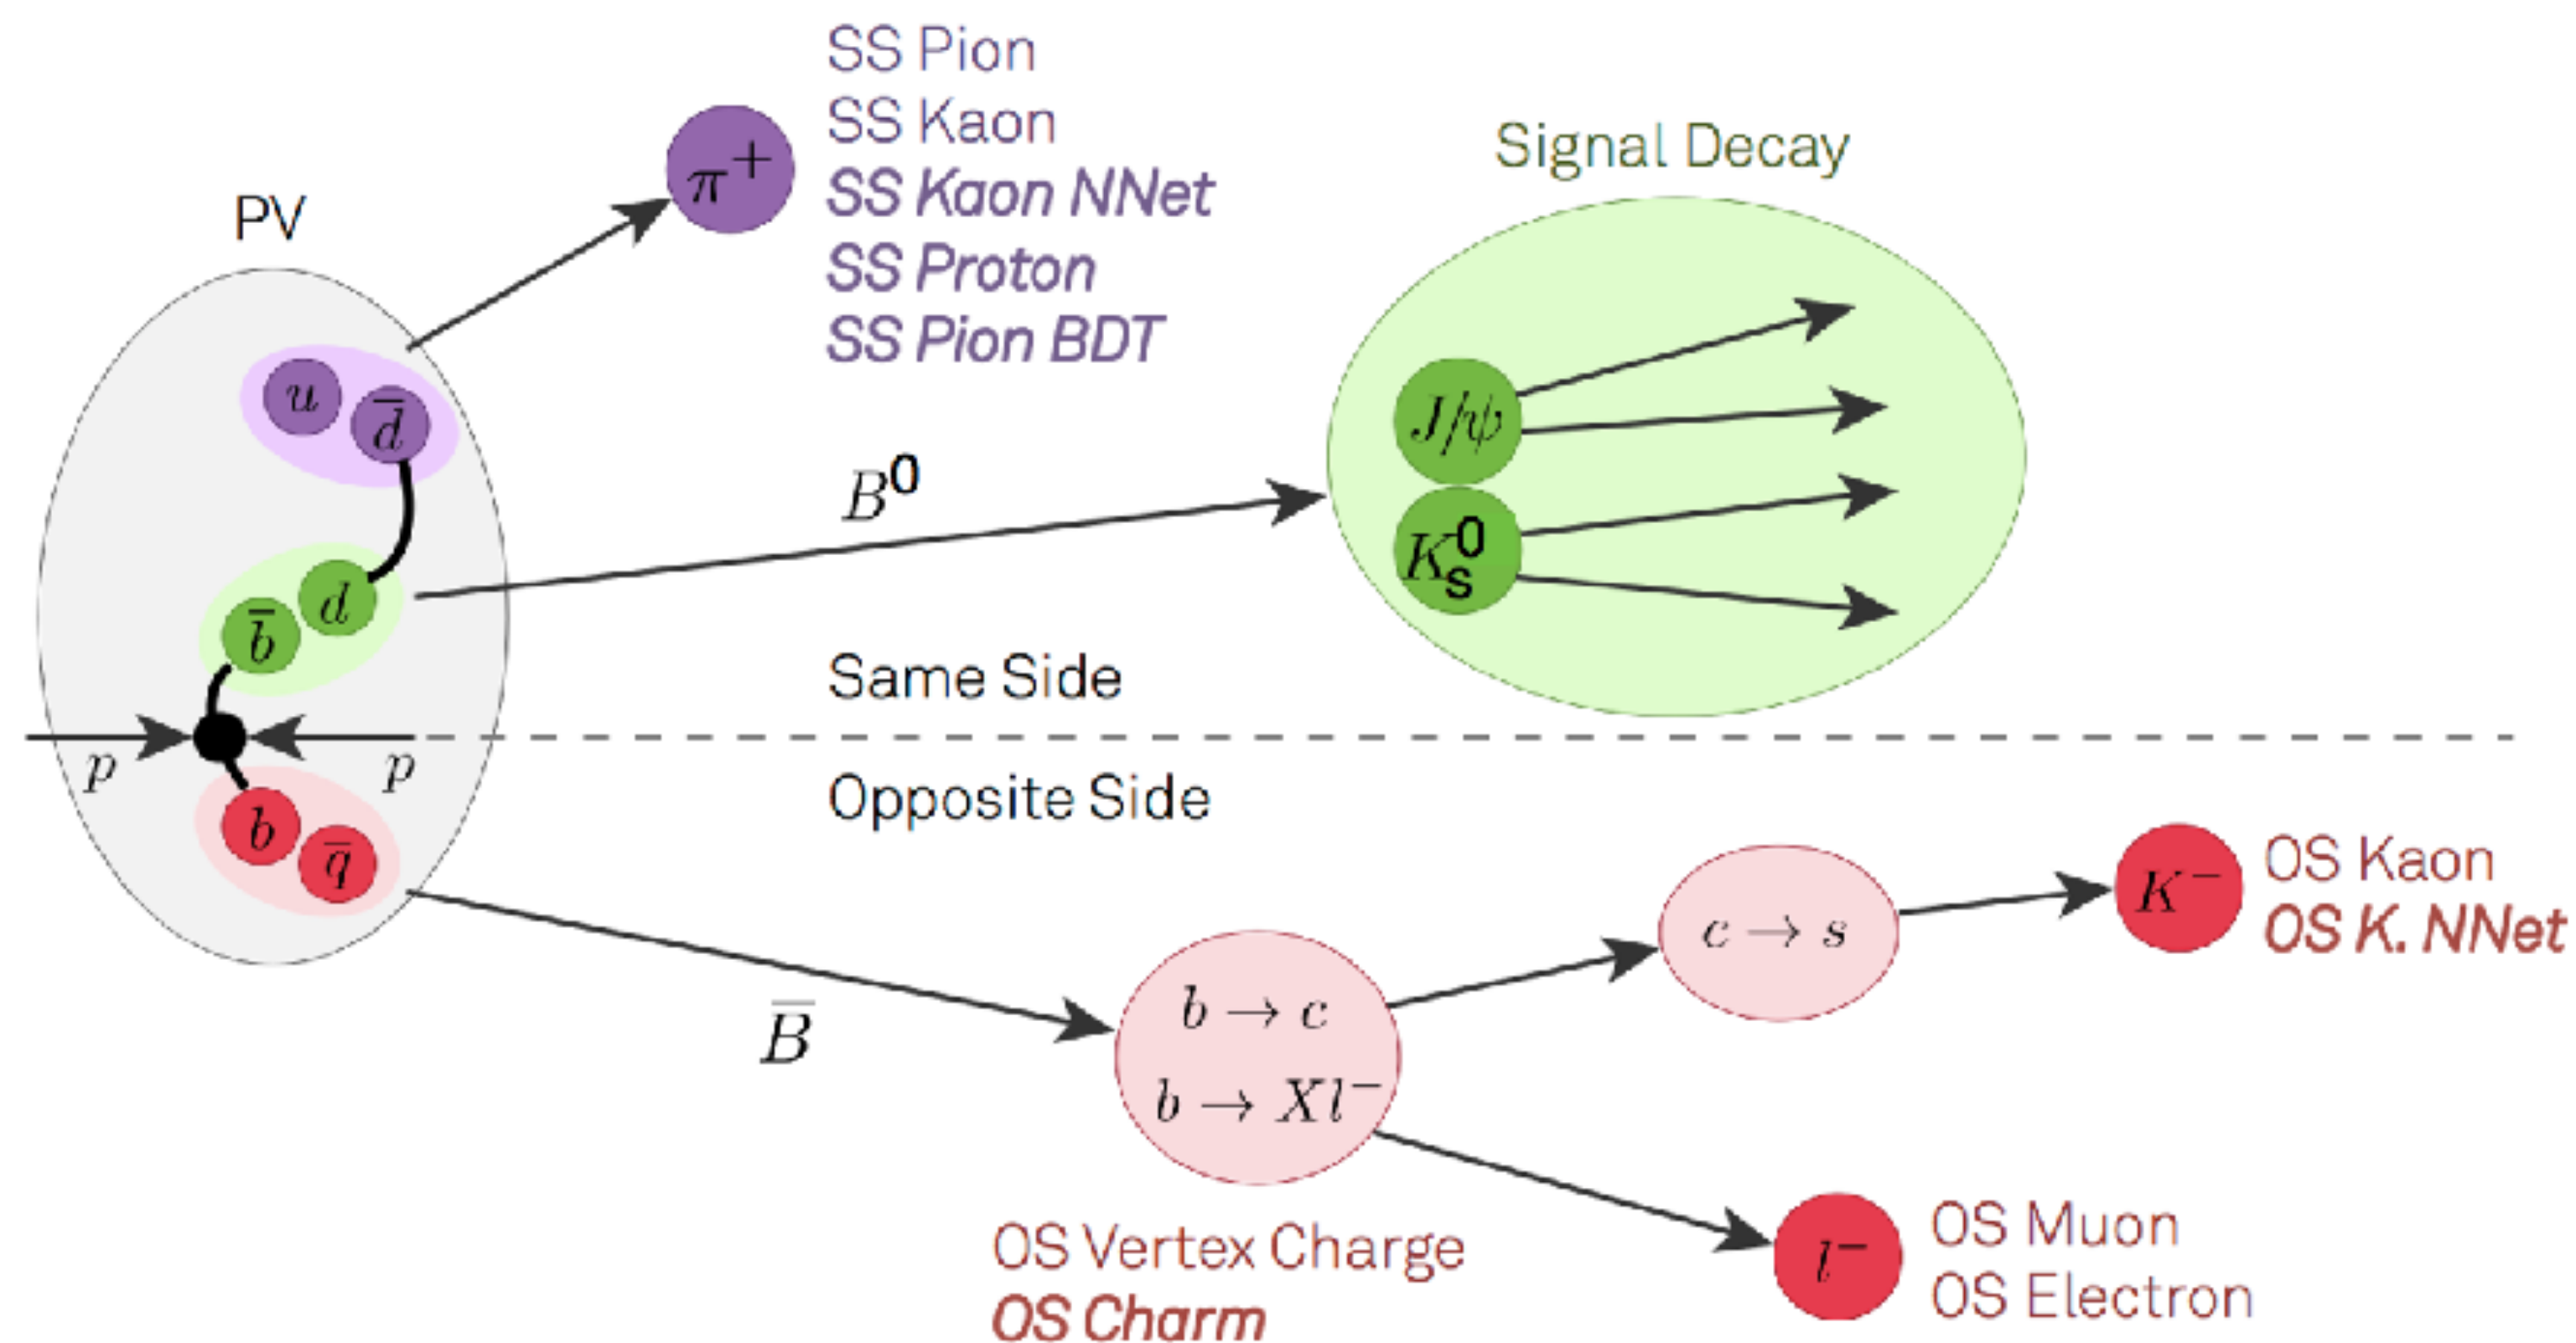
Approaches

- › Ignore negative weights
- › Take $sWeight > 0$ as signal, $sWeight \leq 0$ as bck (with $-sWeight$)
- › Put sample with $weight = P(\text{to be signal})$ as signal and with $weight = P(\text{to be bck})$ as bck
- › Introduce another loss function

B-Tagging



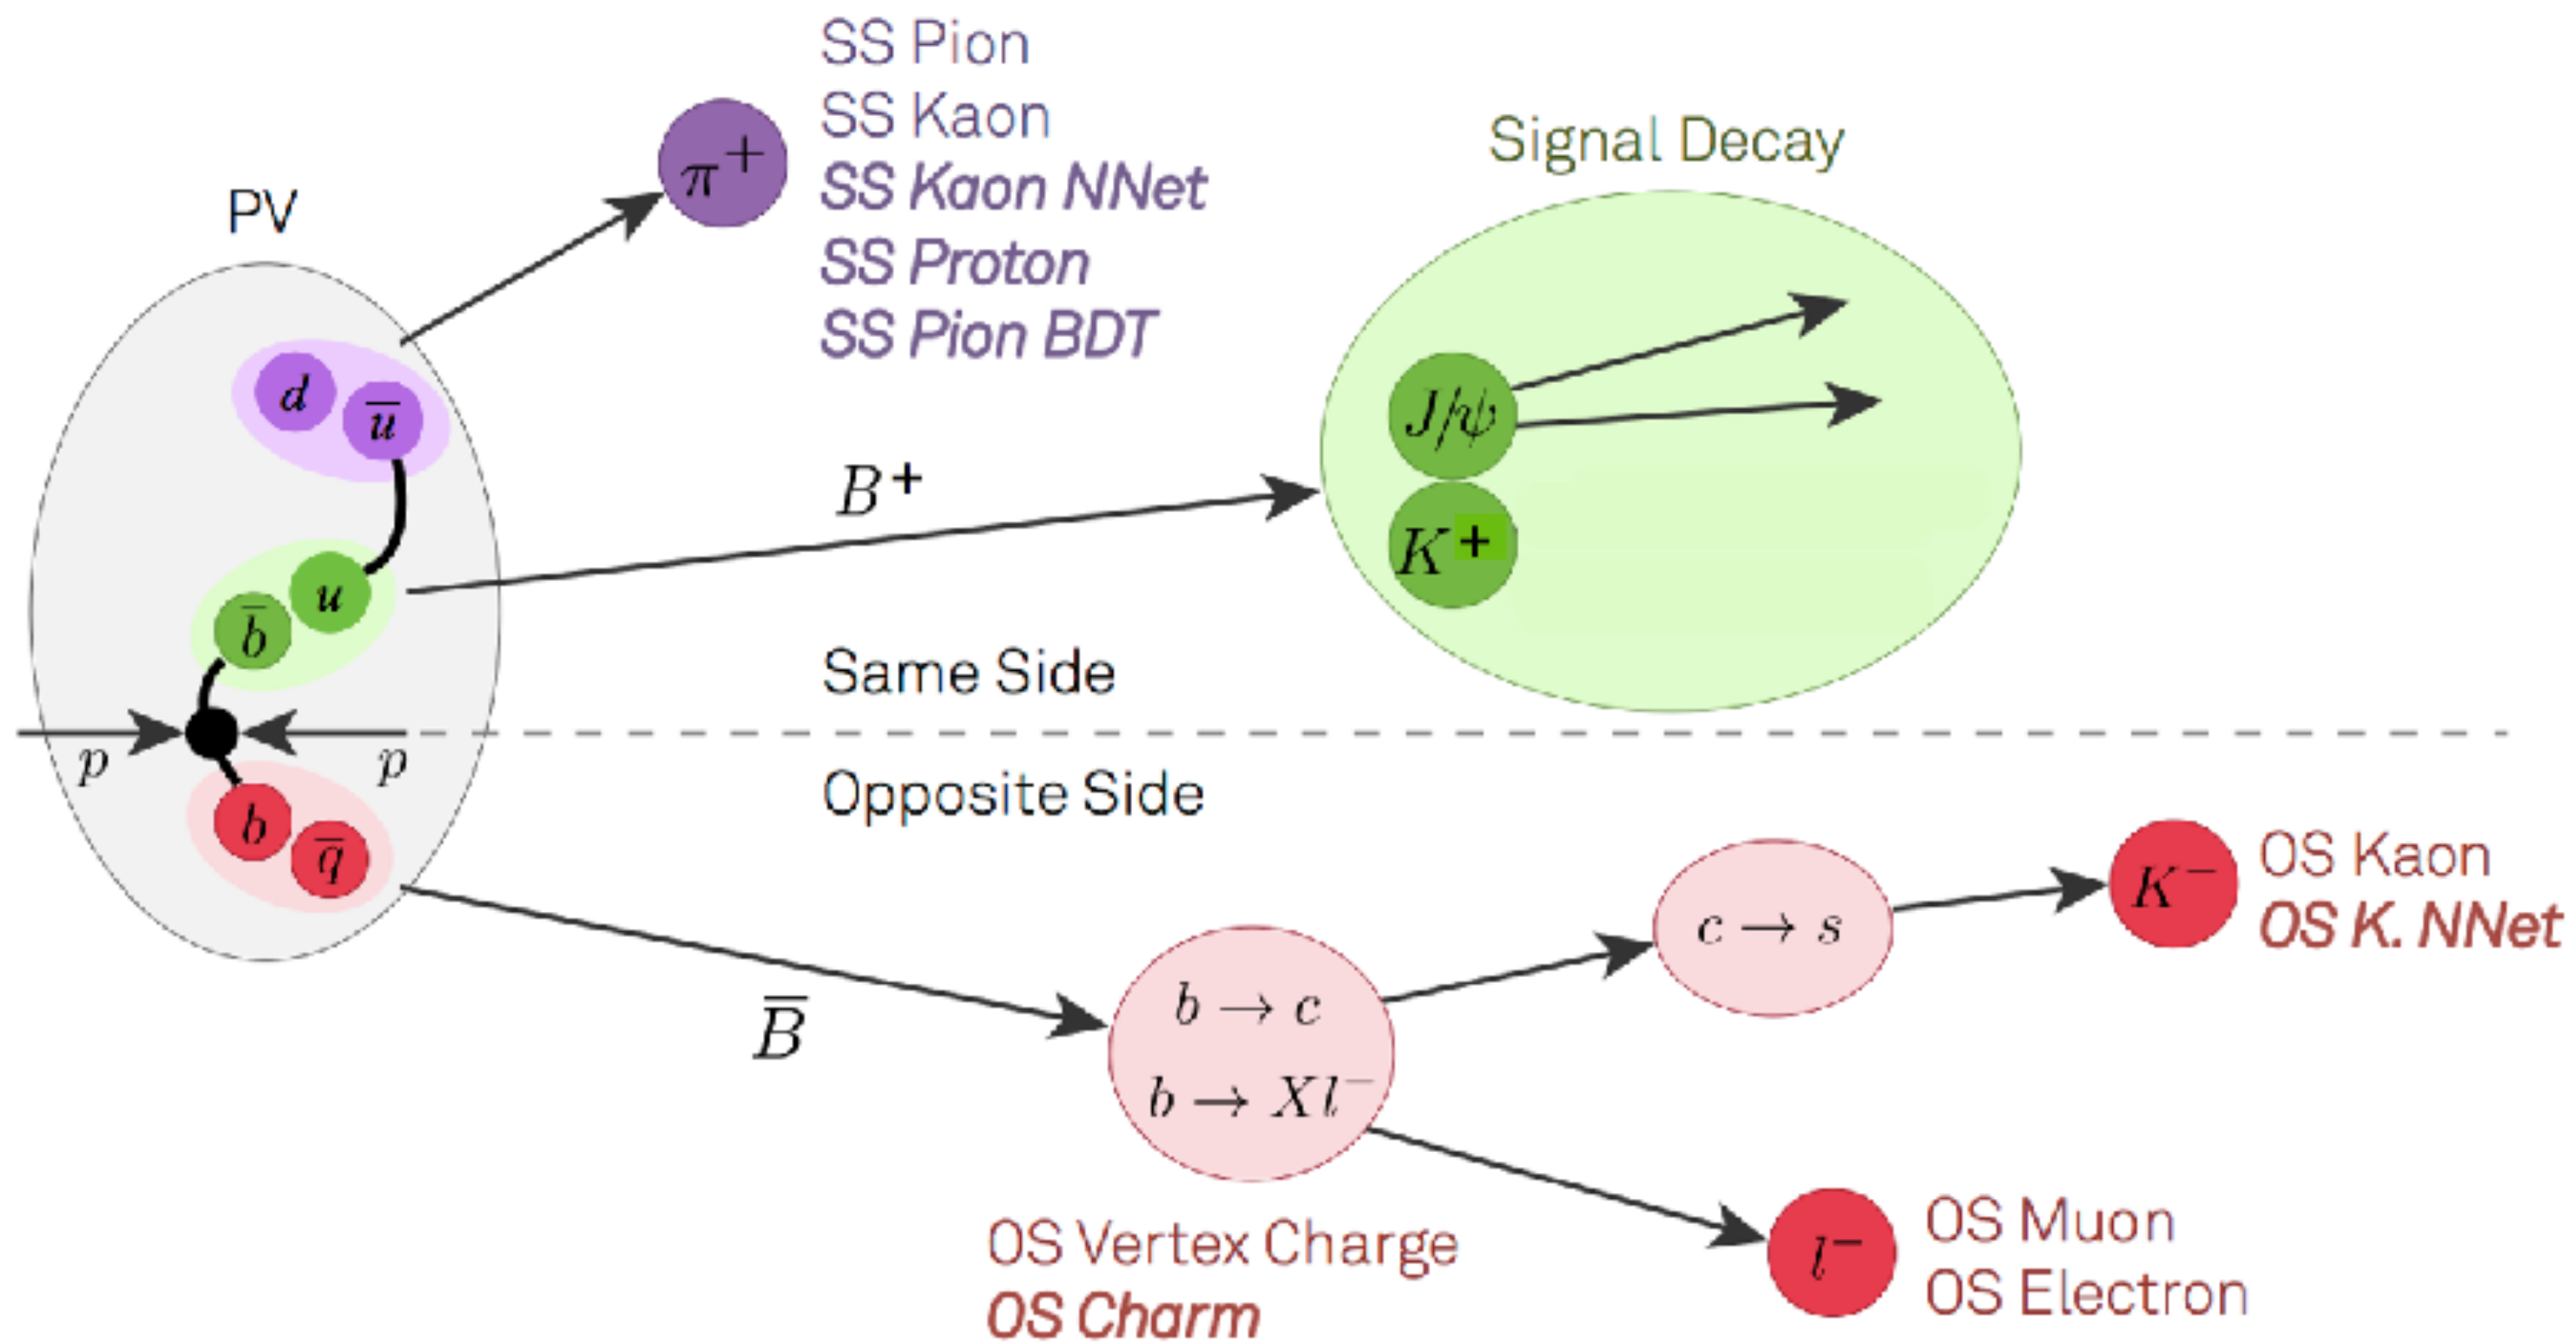
B-Tagging



What is FT?

- › The Flavour Tagging (FT) algorithm determines the flavour of each reconstructed signal B meson at production
- › The B meson consists either of a b or a \bar{b} quark, which defines its flavour
- › The FT algorithm should predict probabilities $P(b)$ and $P(\bar{b})$

B-Tagging (ground truth data)



$$P(\bar{b}) = P(B^+), P(b) = P(B^-);$$

Current tagging system

- › Take $B^\pm \rightarrow J/\psi K^\pm$
- › Apply sPlot technique to real data to extract signal-like data (sWeights)
- › Choose:
 - one tagging track (PID selection and, if necessary, max PT track) for each event
 - one secondary vertex, which produces a tagging particle
- › Train exclusive taggers for SS (same side) tagging particles and OS (opposite side)
- › Target for classifier is ‘right tagged’ label for chosen track/vertex
- › Combine all taggers to one (probabilistic model) to obtain $P(\text{anti-b quark})$

Inclusive probabilistic model

- › Let s_p – charge of track or vertex (+1 or -1), s_B – quark flavour (+1 for b and -1 for \bar{b}), components={tracks, vertices}, then assume that:

$$\begin{aligned} \frac{P(\bar{b})}{P(b)} &= \prod_{\text{components}} \frac{P(\bar{b}|B, \text{component}, s_p)}{P(b|B, \text{component}, s_p)} = \\ &= \prod_{\text{components}} \left(\frac{P(s_B \cdot s_p > 0|B, \text{component})}{P(s_B \cdot s_p < 0|B, \text{component})} \right)^{s_p} \end{aligned}$$

$$flavour = \begin{cases} \bar{b}, & \text{if } P(\bar{b}) \geq P(b) \\ b, & \text{if } P(\bar{b}) < P(b) \end{cases} \quad \omega = \min [P(\bar{b}), P(b)]$$

Why Inclusive?

- › does not depend on the tagging particle type (pion, kaon, electron, muon, proton);
- › is not split into the opposite and same side;
- › combines full available information in the event.

Inclusive training

- › LHCb data samples with reconstructed signal decay $B^+ \rightarrow J/\psi K^+$ or $B^- \rightarrow J/\psi K^-$
- › set of all tracks/vertices for all events form the tracks/vertices datasets (we except for ones forming the reconstructed signal decay)
- › train gradient boosted decision trees (GBDT)
- › target for a classifier is a label that B meson has the same sign as the track/vertex and output is a probability


$$target = \begin{cases} 1, & \text{if } s_B \cdot s_p > 0 \\ 0, & \text{if } s_B \cdot s_p < 0 \end{cases}$$

$$P(s_B \cdot s_p > 0 | B, \text{component})$$

Result

- › ROC AUC score **0.641**
- › ROC AUC for current tagging system **0.566**

ROC curve is computed including untagged events (events which didn't pass preselections; for them probability is set to 0.5)




Classifiers output calibration
to probabilities

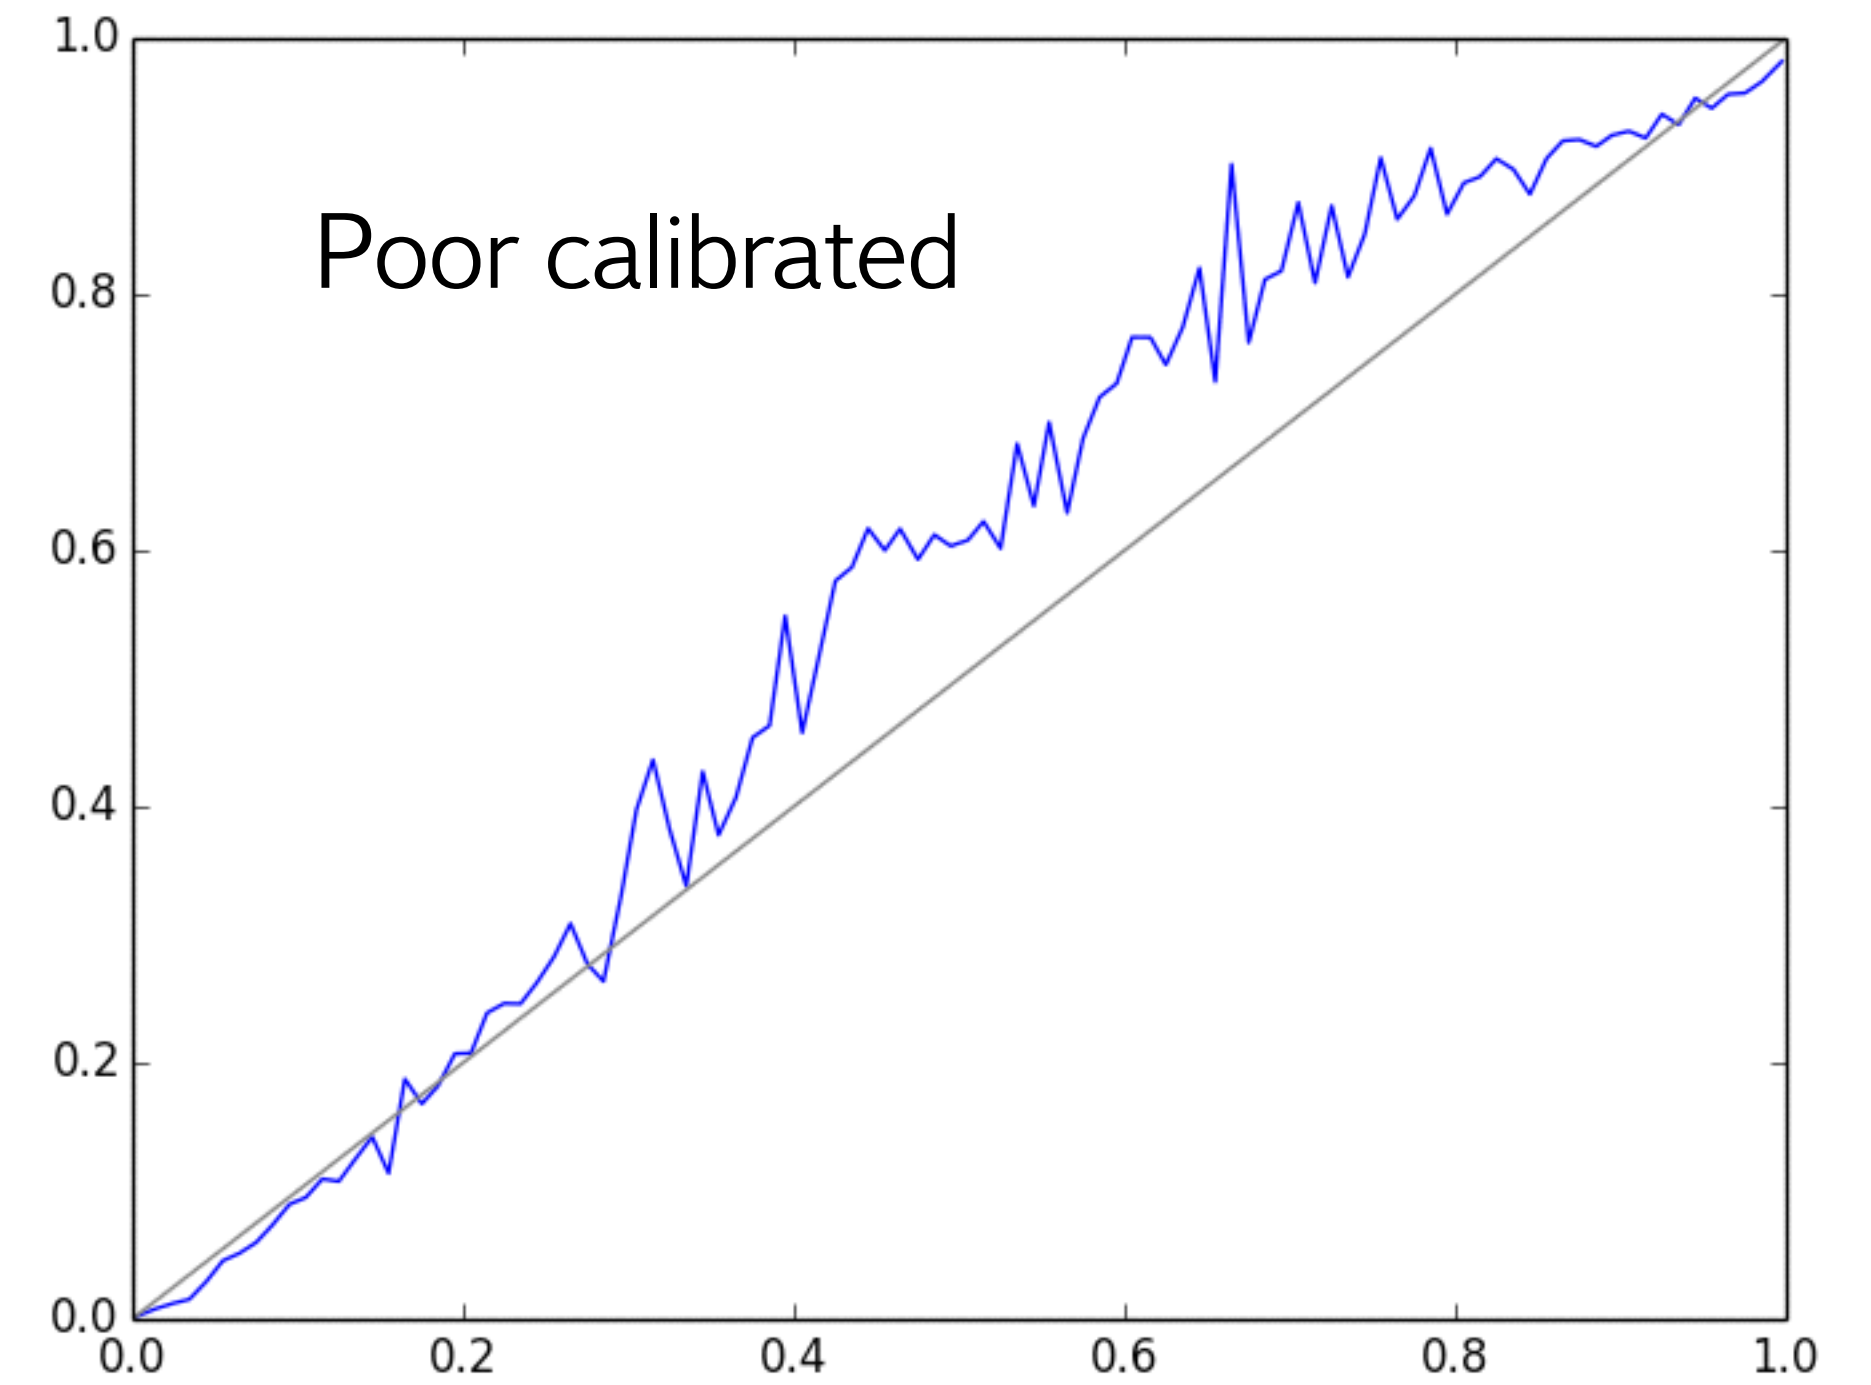
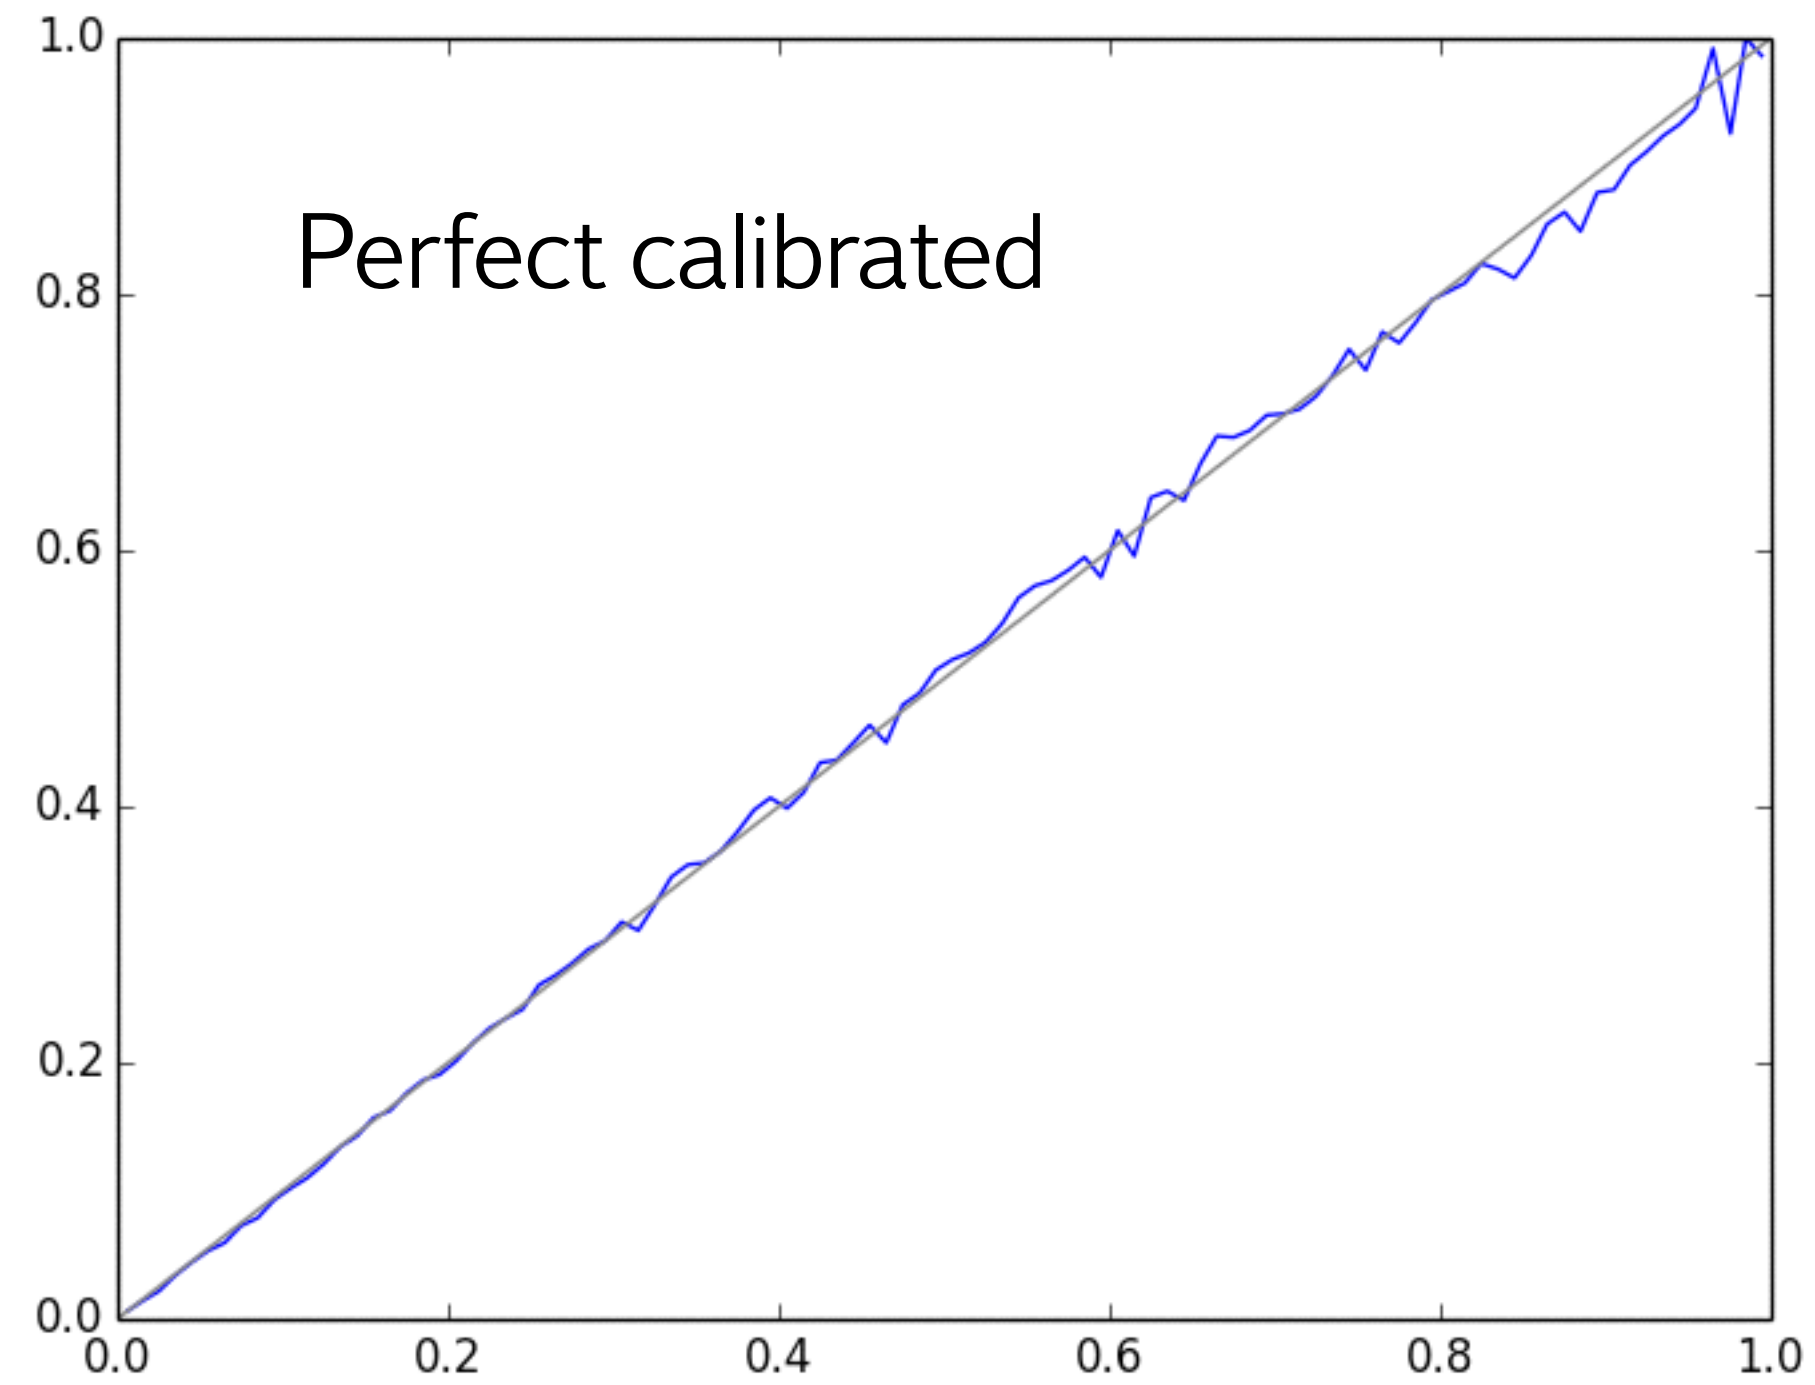


How to check calibration?

- › Enough statistics
- › Divide into many bins output of the classifier (probability to be 1-labeled)
- › Compute in each bin $\#(\text{label}=1) / (\#(\text{label}=1) + \#(\text{label}=0))$
- › Compare it with the mean of predictions in a bin



Examples




Bins approach

› Bins method:


- $\#(1\text{-labeled data in bin}) / \#(\text{data in bin})$
- fit with linear function

› Problems:

- $\#$ bins, their thresholds
- why linear fit?



Platt's scaling (logistic regression)





```
from sklearn.linear_model import LogisticRegression as LR
```

```
lr = LR()
```

```
lr.fit(p_train.reshape( -1, 1 ), y_train )
```

```
p_calibrated = lr.predict_proba( p_test.reshape( -1, 1 ))[:,1]
```

Isotonic regression



Sklearn interface

Summary



- › Use isotonic calibration
- › Use holdout to check your calibration rule

Calibration in B-tagging





Symmetry problem

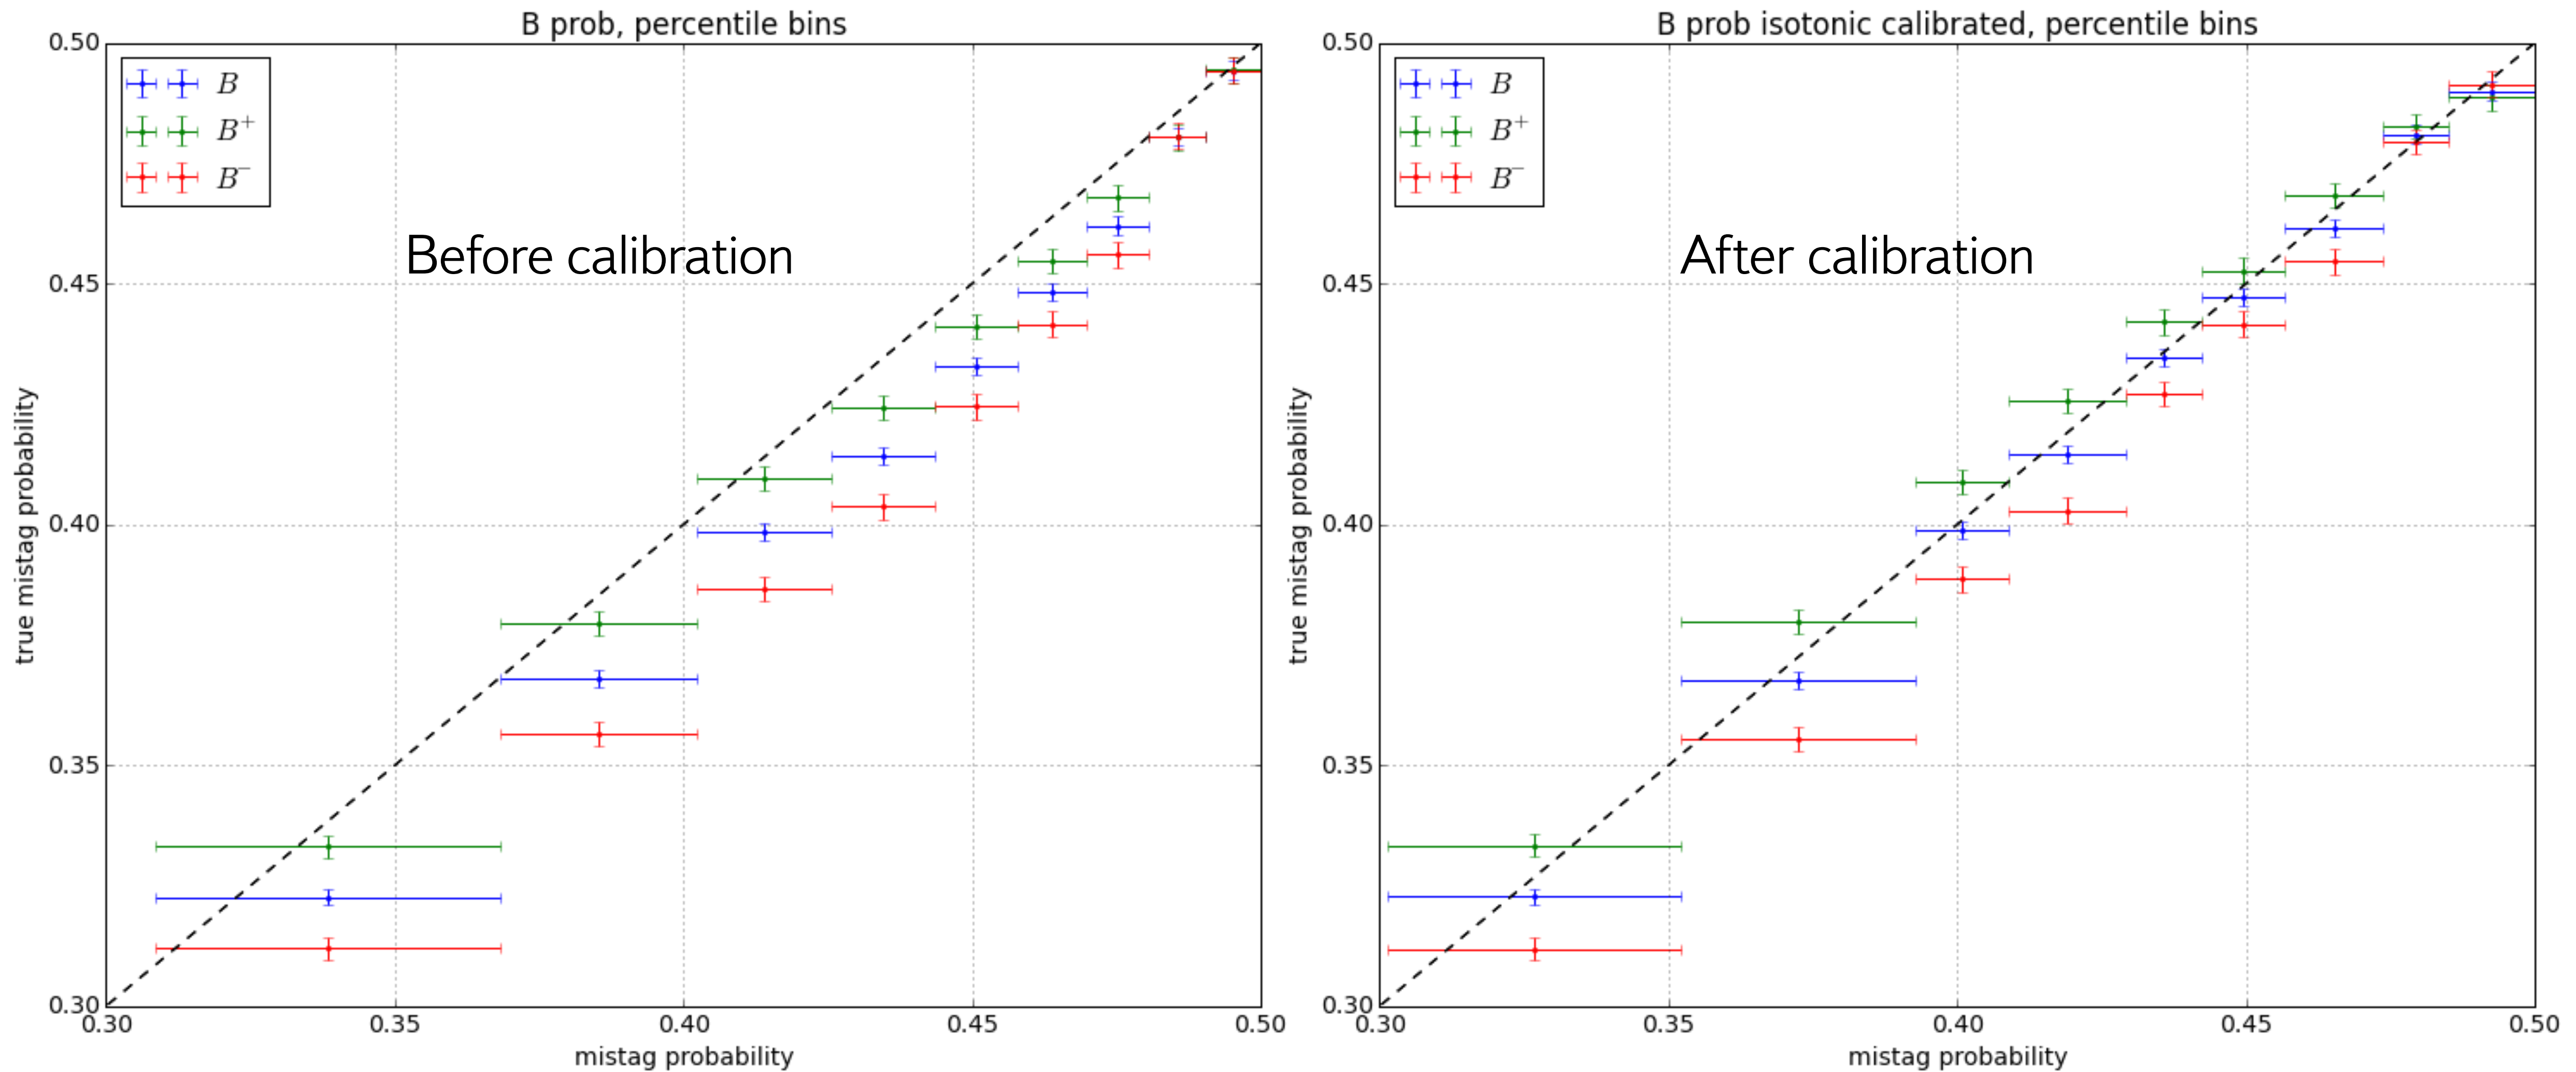
- › symmetric isotonic regression (add inverse labels with inverse probability)
- › add random noise after calibration for stability (to avoid clipped predictions).



Check calibration (with B-symmetry)



Check calibration (with B-symmetry)



Conclusions



Summary & Tricks

- › Measure ROC curve for events!
- › Random forest to remove outliers
- › sPlot technique and ML on sPlot data
- › Try to combine all information for event
- › Calibration aspects: isotonic regression (or its symmetric case with adding noise)
- › Percentile bins (if we need to use binning)

REP

- | Python-based (numpy, pandas, ...), Jupyter-friendly
- | Unified scikit-learn-like API to many ML packages (Sklearn, XGBoost, uBoost, TMVA, Theanets, ...)
- | Meta-algorithms pipelines («REP lego»)
- | Configurable interactive reporting & visualization to ensure model quality (e.g. check for overfitting)
- | Pluggable quality metrics
- | Parallelized training of classifiers & grid search (IPython parallel)
- | Available at [Github](#)

References

- › [LHCb Topological Trigger Reoptimization for Run 2 \(repository\)](#)
- › [LHCb Topological Trigger Reoptimization for Run 2 \(paper\)](#)
- › [The HLT inclusive B triggers](#)
- › [sPlot](#)
- › [Blogpost about sPlot \(simplified explanation\)](#)
- › [Development of "same side" flavour tagging algorithms for measurements of flavour oscillations and CP violation](#)
- › [Blogpost about classifier's output calibration to probability](#)
- › [REP platform](#)

Thanks for attention

Contacts

Likhomanenko Tatiana
researcher-developer



antares@yandex-team.ru, tatiana.likhomanenko@cern.ch