

# Popular ensemble methods

Victor Kitov

Yandex School of Data Analysis



# Table of contents

- 1 Popular ensemble methods
  - Bagging and random forest
  - Boosting

- 1 Popular ensemble methods
  - Bagging and random forest
  - Boosting

# Bagging

- Random selection of
  - samples (with replacement)
  - features (without replacement)
- During bootstrap approximately  $1 - 1/e \approx 2/3$  samples are retained and  $1/e \approx 1/3$  samples left out for large training sets.

# Random forests

**Input:** training dataset  $TDS = \{(x_i, y_i), 1 = 1, 2, \dots, n\}$ ; the number of trees  $B$  and the size of feature subsets  $m$ .

- ❶ for  $b = 1, 2, \dots, B$ :
  - ❶ generate random training dataset  $TDS^b$  of size  $n$  by sampling  $(x_i, y_i)$  pairs from  $TDS$  with replacement.
  - ❷ build a tree using  $TDS^b$  training dataset with feature selection for each node from random subset of features of size  $m$  (generated individually for each node).
- ❷ Evaluate the quality by assigning output to  $x_i, i = 1, 2, \dots, n$  using majority vote (classification) or averaging (regression) among trees with  $b \in \{b : (x_i, y_i) \notin T^b\}$

**Output:**  $B$  trees. Classification is done using majority vote and regression using averaging of  $B$  outputs.

# Comments

- Random forests use random selection on both samples and features
- Left out samples are used for evaluation of model performance.
- Less interpretable than individual trees
- Pro: Parallel implementation
- Contra: different trees are not targeted to correct mistakes of each other

- 1 Popular ensemble methods
  - Bagging and random forest
  - **Boosting**

## Forward stagewise additive modeling

**Input:** training dataset  $(x_i, y_i)$ ,  $i = 1, 2, \dots, n$ ; loss function  $L(f, y)$ , general form of additive classifier  $h(x, \gamma)$  (dependent from parameter  $\gamma$ ) and the number  $M$  of successive additive approximations.

- ① Fit initial approximation  $f^0(x)$  (might be taken  $f^0(x) \equiv 0$ )
- ② For  $m = 1, 2, \dots, M$ :
  - ① find next best classifier

$$(c_m, \gamma_m) = \arg \min \sum_{i=1}^n L(f_{m-1}(x_i) + c_m h(x_i, \gamma_m), y_i)$$

- ② set

$$f_m(x) = f_{m-1}(x) + c_m h(x, \gamma_m)$$

**Output:** approximation function  $f^M(x) = f^0(x) + \sum_{j=1}^M c_j h(x, \gamma_j)$   
Adaboost algorithm is obtained for  $L(y, f(x)) = e^{-yf(x)}$



# Adaboost (discrete version)

**Assumptions:** loss function  $L(y, f(x)) = e^{-yf(x)}$

**Input:** training dataset  $(x_i, y_i)$ ,  $i = 1, 2, \dots, n$ ; number of additive weak classifiers  $M$ , a family of weak classifiers  $h(x) \in \{+1, -1\}$ , trainable on weighted datasets.

- 1 Initialize observation weights  $w_i = 1/n$ ,  $i = 1, 2, \dots, n$ .
- 2 for  $m = 1, 2, \dots, M$ :
  - 1 fit  $h^m(x)$  to training data using weights  $w_i$
  - 2 compute weighted misclassification rate:

$$E_m = \frac{\sum_{i=1}^n w_i \mathbb{I}[h^m(x) \neq y_i]}{\sum_{i=1}^n w_i}$$

- 3 if  $E_m > 0.5$  or  $E_m = 0$ : terminate procedure.
- 4 compute  $\alpha_m = \ln((1 - E_m)/E_m)$
- 5 increase all weights, where misclassification with  $h^m(x)$  was made:

$$w_i \leftarrow w_i e^{\alpha_m}, i \in \{i : h^m(x_i) \neq y_i\}$$

**Output:** composite classifier  $f(x) = \text{sign} \left( \sum_{m=1}^M \alpha_m h^m(x) \right)$

## Adaboost derivation

Set initial approximation  $f^0(x) \equiv 0$ .

Apply forward stagewise algorithm for  $m = 1, 2, \dots, M$ :

$$\begin{aligned}(c_m, h^m) &= \arg \min_{c_m, h^m} \sum_{i=1}^n L(f_{m-1}(x_i) + c_m h^m(x), y_i) \\&= \arg \min_{c_m, h^m} \sum_{i=1}^n e^{-y_i f_{m-1}(x_i)} e^{-c_m y_i h^m(x)} \\&= \arg \min_{c_m, h^m} \sum_{i=1}^n w_i^m e^{-c_m y_i h^m(x_i)}, \quad w_i^m = e^{-y_i f_{m-1}(x_i)}\end{aligned}$$

Since  $c_m \geq 0$  and  $y_i h^m(x_i) \in \{-1, +1\}$  minimum with respect to  $h^m(x)$  is attained at

$$h^m(x_i) = \arg \min_h \sum_{i=1}^n w_i^m \mathbb{I}[h(x_i) \neq y_i]$$

# Adaboost derivation

Denote  $F(c_m) = \sum_{i=1}^n w_i^m \exp(-c_m y_i h^m(x_i))$ . Then

$$\frac{\partial F(c_m)}{\partial c_m} = - \sum_{i=1}^n w_i^m e^{-c_m y_i h^m(x_i)} y_i h^m(x_i) = 0$$

$$- \sum_{i: h^m(x_i) = y_i} w_i^m e^{-c_m} + \sum_{i: h^m(x_i) \neq y_i} w_i^m e^{c_m} = 0$$

$$e^{2c_m} = \frac{\sum_{i: h^m(x_i) = y_i} w_i^m}{\sum_{i: h^m(x_i) \neq y_i} w_i^m}$$

$$c_m = \frac{1}{2} \ln \frac{\left( \sum_{i: h^m(x_i) = y_i} w_i^m \right) / \left( \sum_{i=1}^n w_i^m \right)}{\left( \sum_{i: h^m(x_i) \neq y_i} w_i^m \right) / \left( \sum_{i=1}^n w_i^m \right)} = \frac{1}{2} \ln \frac{1 - E_m}{E_m} = \frac{1}{2} \alpha_m,$$

$$E_m = \frac{\sum_{i=1}^n w_i^m \mathbb{I}[h^m(x_i) \neq y_i]}{\sum_{i=1}^n w_i^m}$$

# Adaboost derivation

Weights recalculation:

$$w_i^{m+1} \stackrel{df}{=} e^{-y_i f_m(x_i)} = e^{-y_i f_{m-1}(x_i)} e^{-y_i c_m h^m(x_i)}$$

Noting that  $-y_i h^m(x_i) = 2\mathbb{I}[h^m(x_i) \neq y_i] - 1$ , we can rewrite:

$$\begin{aligned} w_i^{m+1} &= e^{-y_i f_{m-1}(x_i)} e^{c_m (2\mathbb{I}[h^m(x_i) \neq y_i] - 1)} = \\ &= w_i^m e^{2c_m \mathbb{I}[h^m(x_i) \neq y_i]} e^{-c_m} \propto w_i^m e^{2c_m \mathbb{I}[h^m(x_i) \neq y_i]} = w_i^m e^{\alpha_m \mathbb{I}[h^m(x_i) \neq y_i]} \end{aligned}$$

We used the property that classification result is not affected by multiplication of all weights by constant.

# Gradient boosting

- For general loss function  $L$  forward stagewise algorithm can be solved explicitly in rare cases. In general gradient boosting is applied.
- Gradient boosting is analogous to steepest descent:
  - function approximation is composed of sums of approximations, each of which approximates  $\partial L / \partial f$ .

# Gradient boosting

**Input:** training dataset  $(x_i, y_i)$ ,  $i = 1, 2, \dots, n$ ; loss function  $L(f, y)$  and the number  $M$  of successive additive approximations.

- 1 Fit initial approximation  $f^0(x)$  (might be taken  $f^0(x) \equiv 0$ )

# Gradient boosting

**Input:** training dataset  $(x_i, y_i)$ ,  $i = 1, 2, \dots, n$ ; loss function  $L(f, y)$  and the number  $M$  of successive additive approximations.

- 1 Fit initial approximation  $f^0(x)$  (might be taken  $f^0(x) \equiv 0$ )
- 2 For each step  $m = 1, 2, \dots, M$ :

# Gradient boosting

**Input:** training dataset  $(x_i, y_i)$ ,  $i = 1, 2, \dots, n$ ; loss function  $L(f, y)$  and the number  $M$  of successive additive approximations.

- 1 Fit initial approximation  $f^0(x)$  (might be taken  $f^0(x) \equiv 0$ )
- 2 For each step  $m = 1, 2, \dots, M$ :
  - 1 calculate derivatives  $z_i = -\frac{\partial L(r, y)}{\partial r} \Big|_{r=f^{m-1}(x)}$



# Gradient boosting

**Input:** training dataset  $(x_i, y_i)$ ,  $i = 1, 2, \dots, n$ ; loss function  $L(f, y)$  and the number  $M$  of successive additive approximations.

- 1 Fit initial approximation  $f^0(x)$  (might be taken  $f^0(x) \equiv 0$ )
- 2 For each step  $m = 1, 2, \dots, M$ :
  - 1 calculate derivatives  $z_i = -\frac{\partial L(r, y)}{\partial r} \Big|_{r=f^{m-1}(x)}$
  - 2 fit  $h^m$  on  $\{(x_i, z_i)\}_{i=1}^n$  with some loss function.

# Gradient boosting

**Input:** training dataset  $(x_i, y_i)$ ,  $i = 1, 2, \dots, n$ ; loss function  $L(f, y)$  and the number  $M$  of successive additive approximations.

- ① Fit initial approximation  $f^0(x)$  (might be taken  $f^0(x) \equiv 0$ )
- ② For each step  $m = 1, 2, \dots, M$ :
  - ① calculate derivatives  $z_i = -\frac{\partial L(r, y)}{\partial r} \Big|_{r=f^{m-1}(x)}$
  - ② fit  $h^m$  on  $\{(x_i, z_i)\}_{i=1}^n$  with some loss function.
  - ③ solve univariate optimization problem:

$$\sum_{i=1}^n L(f^{m-1}(x_i) + c_m h^m(x_i), y_i) \rightarrow \min_{c_m \in \mathbb{R}_+}$$

# Gradient boosting

**Input:** training dataset  $(x_i, y_i)$ ,  $i = 1, 2, \dots, n$ ; loss function  $L(f, y)$  and the number  $M$  of successive additive approximations.

- ① Fit initial approximation  $f^0(x)$  (might be taken  $f^0(x) \equiv 0$ )
- ② For each step  $m = 1, 2, \dots, M$ :
  - ① calculate derivatives  $z_i = -\frac{\partial L(r, y)}{\partial r} \Big|_{r=f^{m-1}(x)}$
  - ② fit  $h^m$  on  $\{(x_i, z_i)\}_{i=1}^n$  with some loss function.
  - ③ solve univariate optimization problem:

$$\sum_{i=1}^n L(f^{m-1}(x_i) + c_m h^m(x_i), y_i) \rightarrow \min_{c_m \in \mathbb{R}_+}$$

- ④ set  $f^m(x) = f^{m-1}(x) + c_m h^m(x)$

# Gradient boosting

**Input:** training dataset  $(x_i, y_i)$ ,  $i = 1, 2, \dots, n$ ; loss function  $L(f, y)$  and the number  $M$  of successive additive approximations.

- ① Fit initial approximation  $f^0(x)$  (might be taken  $f^0(x) \equiv 0$ )
- ② For each step  $m = 1, 2, \dots, M$ :
  - ① calculate derivatives  $z_i = -\frac{\partial L(r, y)}{\partial r} \Big|_{r=f^{m-1}(x)}$
  - ② fit  $h^m$  on  $\{(x_i, z_i)\}_{i=1}^n$  with some loss function.
  - ③ solve univariate optimization problem:

$$\sum_{i=1}^n L(f^{m-1}(x_i) + c_m h^m(x_i), y_i) \rightarrow \min_{c_m \in \mathbb{R}_+}$$

- ④ set  $f^m(x) = f^{m-1}(x) + c_m h^m(x)$

**Output:** approximation function  $f^M(x) = f^0(x) + \sum_{m=1}^M c_m h^m(x)$

# Gradient boosting of trees

**Input:** training dataset  $(x_i, y_i)$ ,  $i = 1, 2, \dots, n$ ; loss function  $L(f, y)$  and the number  $M$  of successive additive approximations.

- 1 Fit constant initial approximation  $f^0(x)$ :

$$f^0(x) = \arg \min_{\gamma} \sum_{i=1}^n L(\gamma, y_i)$$

# Gradient boosting of trees

**Input:** training dataset  $(x_i, y_i)$ ,  $i = 1, 2, \dots, n$ ; loss function  $L(f, y)$  and the number  $M$  of successive additive approximations.

- 1 Fit constant initial approximation  $f^0(x)$ :  
$$f^0(x) = \arg \min_{\gamma} \sum_{i=1}^n L(\gamma, y_i)$$
- 2 For each step  $m = 1, 2, \dots, M$ :

# Gradient boosting of trees

**Input:** training dataset  $(x_i, y_i)$ ,  $i = 1, 2, \dots, n$ ; loss function  $L(f, y)$  and the number  $M$  of successive additive approximations.

- 1 Fit constant initial approximation  $f^0(x)$ :  
$$f^0(x) = \arg \min_{\gamma} \sum_{i=1}^n L(\gamma, y_i)$$
- 2 For each step  $m = 1, 2, \dots, M$ :
  - 1 calculate derivatives  $\mathbf{z}_i = -\frac{\partial L(r, y)}{\partial r} \Big|_{r=f^{m-1}(x)}$

# Gradient boosting of trees

**Input:** training dataset  $(x_i, y_i)$ ,  $i = 1, 2, \dots, n$ ; loss function  $L(f, y)$  and the number  $M$  of successive additive approximations.

- 1 Fit constant initial approximation  $f^0(x)$ :  
$$f^0(x) = \arg \min_{\gamma} \sum_{i=1}^n L(\gamma, y_i)$$
- 2 For each step  $m = 1, 2, \dots, M$ :
  - 1 calculate derivatives  $z_i = -\frac{\partial L(r, y)}{\partial r} \Big|_{r=f^{m-1}(x)}$
  - 2 fit regression tree  $h^m$  on  $\{(x_i, z_i)\}_{i=1}^n$  with some loss function, get leaf regions  $\{R_{jm}\}_{j=1}^{J_m}$ .



# Gradient boosting of trees

**Input:** training dataset  $(x_i, y_i)$ ,  $i = 1, 2, \dots, n$ ; loss function  $L(f, y)$  and the number  $M$  of successive additive approximations.

- ① Fit constant initial approximation  $f^0(x)$ :  
$$f^0(x) = \arg \min_{\gamma} \sum_{i=1}^n L(\gamma, y_i)$$
- ② For each step  $m = 1, 2, \dots, M$ :
  - ① calculate derivatives  $z_i = -\frac{\partial L(r, y)}{\partial r} \Big|_{r=f^{m-1}(x)}$
  - ② fit regression tree  $h^m$  on  $\{(x_i, z_i)\}_{i=1}^n$  with some loss function, get leaf regions  $\{R_{jm}\}_{j=1}^{J_m}$ .
  - ③ for each terminal region  $R_{jm}$ ,  $j = 1, 2, \dots, J_m$  solve univariate optimization problem:

$$\gamma_{jm} = \arg \min_{\gamma} \sum_{x_i \in R_{jm}} L(f^{m-1}(x_i) + \gamma, y_i)$$

# Gradient boosting of trees

**Input:** training dataset  $(x_i, y_i)$ ,  $i = 1, 2, \dots, n$ ; loss function  $L(f, y)$  and the number  $M$  of successive additive approximations.

- ① Fit constant initial approximation  $f^0(x)$ :  

$$f^0(x) = \arg \min_{\gamma} \sum_{i=1}^n L(\gamma, y_i)$$
- ② For each step  $m = 1, 2, \dots, M$ :
  - ① calculate derivatives  $z_i = -\frac{\partial L(r, y)}{\partial r} \Big|_{r=f^{m-1}(x)}$
  - ② fit regression tree  $h^m$  on  $\{(x_i, z_i)\}_{i=1}^n$  with some loss function, get leaf regions  $\{R_{jm}\}_{j=1}^{J_m}$ .
  - ③ for each terminal region  $R_{jm}$ ,  $j = 1, 2, \dots, J_m$  solve univariate optimization problem:

$$\gamma_{jm} = \arg \min_{\gamma} \sum_{x_i \in R_{jm}} L(f^{m-1}(x_i) + \gamma, y_i)$$

- ④ update  $f^m(x) = f^{m-1}(x) + \sum_{j=1}^{J_m} \gamma_{jm} \mathbb{I}[x \in R_{jm}]$

# Gradient boosting of trees

**Input:** training dataset  $(x_i, y_i)$ ,  $i = 1, 2, \dots, n$ ; loss function  $L(f, y)$  and the number  $M$  of successive additive approximations.

- ① Fit constant initial approximation  $f^0(x)$ :  
$$f^0(x) = \arg \min_{\gamma} \sum_{i=1}^n L(\gamma, y_i)$$
- ② For each step  $m = 1, 2, \dots, M$ :
  - ① calculate derivatives  $z_i = -\frac{\partial L(r, y)}{\partial r} \Big|_{r=f^{m-1}(x)}$
  - ② fit regression tree  $h^m$  on  $\{(x_i, z_i)\}_{i=1}^n$  with some loss function, get leaf regions  $\{R_{jm}\}_{j=1}^{J_m}$ .
  - ③ for each terminal region  $R_{jm}$ ,  $j = 1, 2, \dots, J_m$  solve univariate optimization problem:

$$\gamma_{jm} = \arg \min_{\gamma} \sum_{x_i \in R_{jm}} L(f^{m-1}(x_i) + \gamma, y_i)$$

- ④ update  $f^m(x) = f^{m-1}(x) + \sum_{j=1}^{J_m} \gamma_{jm} \mathbb{I}[x \in R_{jm}]$

**Output:** approximation function  $f^M(x)$

## Modification of boosting for trees

- Compared to first method of gradient boosting, boosting of regression trees finds additive coefficients individually for each terminal region  $R_{jm}$ , not globally for the whole classifier  $h^m(x)$ .
- This is done to increase accuracy: forward stagewise algorithm cannot be applied to find  $R_{jm}$ , but it can be applied to find  $\gamma_{jm}$ , because second task is solvable for arbitrary  $L$ .
- Max leaves  $J$ 
  - interaction between no more than  $J - 1$  terms
  - usually  $4 \leq J \leq 8$
  - $M$  controls underfitting-overfitting tradeoff and selected using validation set

# Loss function selection

- Usually  $h$  is fitted to  $z$  using  $(h - z)^2$ , though using absolute loss or Huber function loss makes procedure more robust to outliers.
- $L(\hat{y}, y)$  specification:
  - in regression  $L(\hat{y}, y)$  is set to regression loss function:
    - squared deviation, absolute deviation, Huber loss
    - when  $L(\hat{y}, y) = (\hat{y} - y)^2$  method is called *L2Boost*.
  - in classification  $L(f, y)$  is set to margin loss function:
    - exponential  $L(f, y) = e^{-fy}$  or log-loss  $L(f, y) = \ln(1 + e^{-fy})$ .
    - log-loss optimization with approximate solution from single step of Newton-Raphson method is called *LogitBoost*
    - log-loss optimization yields not only classes, but also class probabilities.

# Shrinkage & subsampling

- Shrinkage of general GB, step (d):

$$f^m(x) = f^{m-1}(x) + \nu c_m h^m(x)$$

- Shrinkage of trees GB, step (d):

$$f^m(x) = f^{m-1}(x) + \nu \sum_{j=1}^{J_m} \gamma_{jm} \mathbb{I}[x \in R_{jm}]$$

- Comments:

- $\nu \in (0, 1]$
- $\nu \downarrow \implies M \uparrow$

- Subsampling